

TTK4250 Sensor Fusion

Assignment 4

Hand in: *Wednesday 23. September 16.00* on Blackboard.

Tasks are to be solved on paper if you are not told otherwise, and you are supposed to show how you got to a particular answer. It is, however, encouraged to use Python, MATLAB, Maple, etc. to verify your answers. Rottmann's mathematical formula collection is allowed at both the exercises and the exam.

Task 1: *Gaussian mixture reduction*

- In the file `mixture_reduction.py`, implement the Python function `merge`. This should implement equations 6.19–6.21 in the book.
- You are waiting for a friend who is driving to you on a particular route. While you are waiting, you deduce a distribution over how far you believe he has come from a couple of messages from him. The messages were a bit unclear, so you ended up with a univariate Gaussian mixture with three components as the distribution. However, you think it is a bit much with three components and believe that you should get a good enough representation with only two Gaussians. So, you decide to merge two of them by making these two into a new Gaussian that matches the mean and covariance of the mixture of the two. To keep it simple you neglect any time/speed aspect.

In the given Gaussian mixtures, which would you merge by moment matching if you were to merge two components, why would you merge these, and what would the resulting components be?

The script `perform_mixture_reduction.py` is given if you want to visualize the problem and use your solution to a) for calculations.

Hint: Section 6.3 in the book can provide some insight. For merging you can use the fact that $w_1 p_1(x) + w_2 p_2(x) + w_3 p_3(x) = w_1 p_1(x) + (w_2 + w_3) \left(\frac{w_2}{w_2 + w_3} p_2(x) + \frac{w_3}{w_2 + w_3} p_3(x) \right)$. There is not necessarily a one true answer here, and it depends on what you emphasize.

- $w = \{\frac{1}{3}, \frac{1}{3}, \frac{1}{3}\}$, $\mu = \{0, 2, 4.5\}$ and $P = \{1^2, 1^2, 1^2\}$.
- $w = \{\frac{1}{6}, \frac{4}{6}, \frac{1}{6}\}$, $\mu = \{0, 2, 4.5\}$ and $P = \{1^2, 1^2, 1^2\}$.
- $w = \{\frac{1}{3}, \frac{1}{3}, \frac{1}{3}\}$, $\mu = \{0, 2, 4.5\}$ and $P = \{1^2, 1.5^2, 1.5^2\}$.
- $w = \{\frac{1}{3}, \frac{1}{3}, \frac{1}{3}\}$, $\mu = \{0, 0, 2.5\}$ and $P = \{1^2, 1.5^2, 1.5^2\}$.

Task 2: *Measurement likelihood of the interacting multiple model filter and particle filter*

Hint: These should not end up being very complicated.

- Derive the measurement likelihood, $p(z_k | z_{1:k-1})$, of an IMM filter by using of the total probability theorem,

$$p(z_k | z_{1:k-1}) = \sum_{s_k} \int p(z_k | x_k, s_k) p(x_k | s_k, z_{1:k-1}) \Pr(s_k | z_{1:k-1}) dx_k,$$

in terms of the mode likelihood, $\Lambda_k^{s_k}$, and mode probabilities, $\Pr(s_k | z_{1:k-1})$, assuming $p(z_k | x_k, s_k) = p(z_k | x_k)$.

Hint: Use terms and equations from section 6.4.

- (b) Assume that the PF state distribution is $p(x_k|z_{1:k-1}) \approx \sum_{i=1}^N w_k^i \delta(x_k - x_k^i)$, where $\delta(x)$ is the Dirac delta function, w_k^i are the normalized weights and x_k^i are the particles. Derive the measurement likelihood, $p(z_k|z_{1:k-1})$, of this PF.

Hint: You can use the total probability theorem to be able to use the given approximate PF state distribution.

Task 3: *Implement an IMM class*

Use the Python skeleton for an IMM class that is available at Blackboard. The function and class definitions has to be as given, but you are free to change any other prewritten code, if you want. You have to implement steps 1 through 4 of the IMM algorithm as given in the book, in addition to a 5th estimation step. The 5th step consists of calculating the overall mean and covariance of the state. Step 3, mode matched filtering, will further be divided into a mode matched prediction and mode matched update step. You will then combine these steps into the higher level functions of predict and update.

Hint: For step 2 and 5 you can use the function `reduceGaussianMixture` you made in task 1.

- (a) Implement step 1 in the function `predict` which should return the predicted mode probability and the mixing probabilities based on the transition matrix and the previous mode probabilities.

This is simplified by completing the file `discretebayes.py` and making the function

- (b) Implement step 2 in the function `update` which should return the mixed mean and covariance based on the mixing probabilities and the previous means and covariances.

To make this generalize and reuse code for the upcoming PDA you can implement the function to do the work in the EKF class

- (c) Implement the prediction part of step 3 in the function `predict` which should output the EKF predictions according to the mode.
- (d) Combine the above into a overall predict function
- (e) Implement the update part of step 3 in the function `update` which should update the mean and covariance along with calculation of the measurement log likelihood, $\log(\Lambda_k^{s_k})$, according to the mode.
- (f) Implement step 4 in `update` which should update the mode probabilities and calculate the overall measurement log likelihood based on the prior mode probabilities and the mode measurement log likelihood.

For generalization and for use with PDA you might want to implement

It might be more numerically stable to do this in logarithms and only exponentiate to get probabilities in the end, but not strictly necessary. However, a function `logsumexp`¹ has been imported for you.

- (g) Combine mode matched update and probability update into an overall update function
- (h) Implement the function that gives the state estimate of the IMM, in terms of a mean and covariance, in the function `estimate`
- (i) Take a look at the functions `NISes` and `NEESes` see that you understand what it is doing. This function might come in handy for tuning the filter. Also have a look at `estimate_sequence`.

Task 4: *Tune an IMM*

You are to tune an IMM consisting of a CV and CT model with position measurement on the data in the given ".mat"-file. The data has been created by simulating these two models with deterministic switching times and setting a new turn rate at the same time. You are, however, to estimate these switching times

¹Calculates the logarithm of a sum of exponentiated terms in an array in a more numerically stable way than the naive approach

along with the trajectory by tuning both filters along with the switching probabilities. The initial state is sampled randomly, and the first measurement is from this initial position. A script can be found on Blackboard to get you started.

Both the CV model and CT models can be found on Blackboard. The CV model is the same as the one you made in the previous assignment except for having added functionality as zero padding to be able to combine it with the 5 states of the CV model. The EKF module is the one you were supposed to implement in assignment 2.

You can try to use NIS and/or NEES, but an optimally tuned filter will not necessarily be within the confidence bounds given by the χ^2 distribution, although it should at least be close. Quantities of interest are low RMSE/MSE, low peak deviation, and good maneuver start/end time estimates.

Hint: The models in IMM must have sufficiently different noise properties if one are to estimate maneuver times. If the models are sufficiently different, the transition probabilities determine how “willing” the filter is to switch to a new model, which gives a tradeoff between maneuver detection and precision. The IMM is a suboptimal filter so the tuning process might have to compensate for that (the process with much noise will inflate the covariance of the process with low noise etc. when mixing).

- (a) Perform a crude tuning of a CV model to the data by getting a consistent NIS, without calculating errors to the ground truth. Then do the same for the CT model. Briefly discuss.

Hint: This is very similar to what you did in assignment 2.

- (b) Tune the IMM without comparing to ground truth. Briefly describe your tuning process and why you chose the parameters. Specifically mention how the differences in parameters are changed from having single filters.

Note: Again, do not spend several hours on this. You should get a feel for the algorithm and see what ambiguity it can give to not have ground truth.

- (c) Now tune by comparing to the ground truth. Did you have to change your parameters? Was it easier than without ground truth?

Task 5: *Implement a SIR particle filter for a pendulum*

We want to estimate the position of a pendulum influenced by some random disturbance, with a suboptimal measuring device.

The pendulum dynamic equation is given by $\ddot{\theta} = \frac{g}{l} \sin(\theta) - d\dot{\theta} + \alpha$, where g is the gravitational acceleration, l is the length from the center of rotation to the pendulum, and d is the dampening coefficient. The pendulum is otherwise assumed to be a point mass attached to a massless rod. α represents other forces that acts on the pendulum in the form of a stochastic angular acceleration process. θ is 0 when the rod is at the bottom.

A measuring device is set up L_d below the rotation point, and L_l to the left so that $z_k = h(\theta_k) = \sqrt{(L_d - \cos(\theta_k))^2 + (\sin(\theta_k) - L_l)^2} + w_k$, where $w_k \sim \text{triangle}(E[w_k], a)$, a symmetric triangle distribution with width $2a$, height $\frac{1}{a}$ and peak at $E[w_k]$.

A skeleton that creates data and provides some plot aid and so on can be found on Blackboard.

- (a) Finish the particle filter. To finish it you need to implement particle prediction, weight update and resampling (algorithm 3). How does the filter perform?
- (b) Do not resample the simulated pendulum, but vary L_l and generate new measurements to see how it changes the estimate. Does it remove any problems in the estimation you found above?
- (c) When would you choose to try a PF instead of a EKF and why?